

How I Bought an AI Computer

And Accidentally Became an Open Source Maintainer Again

A Six Month Review of My Strix Halo Rig

linux | ROCm | Local AI | Home Assistant | Coding

winmutt | June 2026

github.com/winmutt

What We're Covering

- The impulse purchase that changed everything
- ■ Hardware reality check (spoiler: it's complicated)
- Linux + ROCm = Pain (but beautiful pain)
- Performance: How fast can my AI think?
- I cut the Alexa cord (and replaced it with OSS)
- Vibe-coded an entire workspace system
- From digital to physical: Concrete signs?
- How buying hardware got me back in open source
- What I'd do differently (and what I'd do again)

The 10-Year GitHub Glow-Up

Before Strix Halo: Ghost in the machine



November 2025: Bought a Strix Halo rig

Result: 3.5x increase in open source contributions

Moral of the story: Hardware buys you time (and motivation)

The Impulse Purchase

Why AMD Strix Halo?

November 22, 2025

"BTW I have a corsair 300 AMD 395 max+ enroute. Sold my 3060 rig today."


Before: RTX 3060 with 12GB VRAM

After: 128GB LPDDR5X-8000 unified memory

Translation: My AI stopped choking on context

Meet My New Overlord

Corsair AI Workstation 300



Note: Corsair 300 shares motherboard with other Strix Halo manufacturers (Sixunitied)

The Price Tag That Made My Wallet Cry

Black Friday 2025: \$1,800

Today's equivalent: ~\$2,700

TL;DR: I got a good deal (mostly)

OUT OF STOCK

CORSAIR AI WORKSTATION 300 DESKTOP PC

Designed for efficiency and adaptive performance, it's ideal for AI de
Powered by AMD Ryzen™ AI Max 300 Series—ready for local LLMs

- XDNA 2 NPU with up to 50 TOPS of AI acceleration
- Advanced dual-fan cooling for optimal thermal performance
- Performance Level Selector for optimized power and speed
- Access powerful AI tools with the Corsair AI Software Suite
- 2-Year Warranty

\$2,699.99

SYSTEM CPU

AMD Ryzen AI MAX 385

AMD Ryzen AI MAX+395
OUT OF STOCK

SYSTEM GPU

AMD Radeon 8060S iGPU
OUT OF STOCK

SYSTEM MEMORY

128GB LPDDR5X 8000
OUT OF STOCK

SYSTEM STORAGE

1TB M.2 NVMe SSD
OUT OF STOCK

4TB (2TB + 2TB) M.2 NVMe SSD
OUT OF STOCK

CRSR Stock: Because Why Not?

My investment strategy: Buy hardware, watch stock




The Pitch:

- 128GB unified memory (LPDDR5X-8000)
- Fastest 'nom' VRAM available
- NPU for AI acceleration
- No discrete GPU needed

Hardware Reality: It's Complicated

AMD Strix Halo Architecture



The Plot Thickens:

- 128GB total, only 96GB accessible
- Memory placement: VRAM vs GTT drama
- *Swapping happens anyway (thanks unified memory!)*

My AI Menagerie

Models in Rotation

- Qwen3.6-35B-A3B-MTP-GGUF (the daily driver)
- *Qwen3.5-122B (when I want to wait 40 minutes)*
- gpt-oss-120b (for variety)
- Qwen3-Coder-Next (my coding co-pilot)
- Flux-2-Klein-9B (for art, apparently)
- Whisper-Large-v3-Turbo (because typing is hard)
- *Plus 8 more I forgot about*


Multi-Token Prediction (MTP)

- *MTP = Multi-Token Prediction (not Mixed Tensor Precision!)*
- Predicts multiple tokens at once
- Translation: My AI thinks faster

Opencode: The Portable Web UI I Actually Use

Cline is good for direct coding

Opencode is portable (I can code from my couch)



- Cline: Good for direct coding tasks
- Opencode: Preferred for portable web UI experience
- Opencode makes it portable for mobile/different locations

The Software Stack: A Love-Hate Relationship

Lemonade vs Ollama

Feature	Lemonade	Ollama
Backend	llama.cpp + ROCm	Various
Observability	Good (I can see tokens)	<ParagraphStyle 'CustomNormal'>
NPU Support	Coming soon™	Patch exists
Stability	Breaks after updates	GPU crashes
Multi-model	Yes	Good luck

The Resurgens Moment

December 14, 2025

"Upgraded OS and ollama last night and sad things happened. GPU crashes about 3-4 prompts in"

Bleeding Edge = Cutting My Finger

- Near-bleeding edge kernels for APU support
- AI apps change faster than I can debug
- *Each update = new drama*
- Stability vs features: pick your poison

Phoenix Rising: Fresh install → AMD GPU update → 2x faster

Performance: How Fast Does My AI Think?

Token Processing Reality


Context Size	TPS	Notes
<100k tokens	~40	Actually usable
100k-200k	~14-18	I have time to get coffee
>200k tokens	~8-18	Go read a book
>64k (ingress)	N/A	My AI needs a nap

May 30, 2026 - 200k Token Processing

Input: 193,795 | Output: 5,997 | TTFT: 37.8 min | TPS: 18.65

Translation: I made dinner while waiting

Performance Timeline



Project WWS: I Vibe-Coded an Entire System

Winmutt Work Spaces

github.com/winmutt/wws


May 12, 2026

"Definitely no Athena and its taken months to get there but this is something I entirely vibe coded."

What is WWS?

- Remote workspace provisioning
- KVM/Podman isolated environments
- code-server (VSCode in browser)
- GitHub OAuth + RBAC
- *Months of work, zero Athena*

WWS Dashboard



Cutting the Alexa Cord

LineageOS on Echo: Because I Said So

December 2025

Started exploring XDA Forums for Echo device hacking

The Process:

1. Unlock bootloader via amonet exploit
2. *Create stock backup (don't skip this!)*
3. Install TWRP recovery
4. Flash LineageOS 18.1 (Android 11)
5. Configure ViewAssist for Home Assistant


Philosophy

"Basically anything with a USB port. If there is a port, there is a way."

Alexa Replacement Timeline

Dec 6	Started exploring XDA forums
Dec 7	HA running on Echo devices
Jan 19	HA setup on Echo Show Gen 2
Jan 21	Everything works (except wake word)
Jan 25	Hey Jarvis to the rescue
Jan 29	End-to-end complete
Jan 7, 2026	Issue #4: Echo 8 mic dies after 24h
Feb 2026	Still debugging (it's fine)

Home Assistant on Echo Show



Issue #4: The Echo 8 Mic That Quit

Audio stops working after ~24 hours

Filed

January 7, 2026

Repository

<https://github.com/amazon-oss/releases>

Status

Open (still waiting, it's fine)

The Problem

- Echo 5: Stable for days
- Echo 8: Dies after ~24 hours
- *Root cause: Unknown (yet)*

TL;DR

I participate in open source bugs now. You're welcome.

From AI to Concrete: My PE's Worst Nightmare

The Motivation

December 1, 2025


"I am going to use it to create a 3d printed cast for a silicone mold so I can some desktop signs out of concrete."

"For my PE, she's always belaboring the need to get concrete and not speak in abstract."

The Process

1. AI generates OpenSCAD design
2. Export to STL
3. 3D print mold
4. Create silicone mold
5. Cast concrete
6. *Sell on Etsy (eventually)*

Concrete Sign Mold Design



Lessons Learned (Mostly)

What Worked ✓

- Cline + Qwen3-Coder: Actually produces code
- Prompt Caching: Things hum now
- AMD GPU Libraries: 2x faster (worth the pain)
- Home Assistant: Alexa is dead, long live OSS
- WWS Project: Months of vibe-coding complete

What Didn't ✗


- Memory: 32GB disappeared forever
- Ollama: Still crashes after updates
- NPU: Coming soon™ (always)
- Context >64k: TPS drops like a stone

Key Insights (Take What You Want)

1. Hardware hype $\neq$ reality (APU challenges are real)
2. Software maturity: bleeding edge = bleeding fingers
3. Context > everything (prompt caching is life)
4. AI coding > traditional IDEs (for me, anyway)
5. AI $\rightarrow$ physical objects (concrete signs, apparently)
6. Echo $\rightarrow$ Home Assistant = OSS win
7. Autonomous dev: possible, painful, worth it
8. Buying hardware got me back in open source

The Catalyst Effect

AMD Strix Halo: Catalyst for OSS Reinvigoration




How Buying Hardware Got Me Back in Open Source

The Catalyst Effect


November 2025: Bought Strix Halo rig. Local AI reinvigorated my love for OSS. Result: 3.5x increase in contributions.

2026: The Year I Came Back




Projects I've Actually Contributed To

Project Contributions (Nov 2025 - Jun 2026)
Post-Strix Halo OSS Reinvigoration



Project Contributions (Nov 2025 - Jun 2026)



What's Next (Probably)

Immediate

- NPU support (still waiting)
- Memory addressing (32GB, come back)
- llama.cpp PRs (because why not)

Long-term

- Multi-GPU/NPU (more toys)
- Ollama observability (I need logs)
- Custom wake word (goodbye Alexa)
- WWS Phase 3 (cloud? maybe)

Questions? (I Have Answers, Mostly)

1. Is Strix Halo production-ready? (It's complicated)
2. When will NPU support mature? (Coming soon™)
3. How do we solve 128GB addressing? (We don't)
4. Best backend for multi-model? (Lemonade, but...)
5. Is autonomous dev the future? (Ask my AI)

github.com/winmutt

Come find me after - I'll show you the rig